

Final_Spectrogram_Generation

June 30, 2025

```
[1]: import numpy as np
      if not hasattr(np, 'long'):
          np.long = int # fix for deprecated alias
```

```
[4]: import os
      import librosa
      import librosa.display
      import matplotlib.pyplot as plt
```

```
[3]: pip install librosa
```

```
Collecting librosa
  Downloading librosa-0.11.0-py3-none-any.whl.metadata (8.7 kB)
Collecting audioread>=2.1.9 (from librosa)
  Downloading audioread-3.0.1-py3-none-any.whl.metadata (8.4 kB)
Requirement already satisfied: numba>=0.51.0 in /opt/conda/lib/python3.12/site-packages (from librosa) (0.61.2)
Requirement already satisfied: numpy>=1.22.3 in /opt/conda/lib/python3.12/site-packages (from librosa) (2.2.6)
Requirement already satisfied: scipy>=1.6.0 in /opt/conda/lib/python3.12/site-packages (from librosa) (1.15.2)
Requirement already satisfied: scikit-learn>=1.1.0 in /opt/conda/lib/python3.12/site-packages (from librosa) (1.7.0)
Requirement already satisfied: joblib>=1.0 in /opt/conda/lib/python3.12/site-packages (from librosa) (1.5.1)
Requirement already satisfied: decorator>=4.3.0 in /opt/conda/lib/python3.12/site-packages (from librosa) (5.2.1)
Collecting soundfile>=0.12.1 (from librosa)
  Downloading soundfile-0.13.1-py2.py3-none-manylinux_2_28_x86_64.whl.metadata (16 kB)
Collecting pooch>=1.1 (from librosa)
  Downloading pooch-1.8.2-py3-none-any.whl.metadata (10 kB)
Collecting soxr>=0.3.2 (from librosa)
  Downloading soxr-0.5.0.post1-cp312-abi3-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (5.6 kB)
Requirement already satisfied: typing_extensions>=4.1.1 in /opt/conda/lib/python3.12/site-packages (from librosa) (4.14.0)
Requirement already satisfied: lazy_loader>=0.1 in
```

```

/opt/conda/lib/python3.12/site-packages (from librosa) (0.4)
Requirement already satisfied: msgpack>=1.0 in /opt/conda/lib/python3.12/site-
packages (from librosa) (1.1.1)
Requirement already satisfied: packaging in /opt/conda/lib/python3.12/site-
packages (from lazy_loader>=0.1->librosa) (25.0)
Requirement already satisfied: llvmlite<0.45,>=0.44.0dev0 in
/opt/conda/lib/python3.12/site-packages (from numba>=0.51.0->librosa) (0.44.0)
Requirement already satisfied: platformdirs>=2.5.0 in
/opt/conda/lib/python3.12/site-packages (from pooch>=1.1->librosa) (4.3.8)
Requirement already satisfied: requests>=2.19.0 in
/opt/conda/lib/python3.12/site-packages (from pooch>=1.1->librosa) (2.32.4)
Requirement already satisfied: charset_normalizer<4,>=2 in
/opt/conda/lib/python3.12/site-packages (from
requests>=2.19.0->pooch>=1.1->librosa) (3.4.2)
Requirement already satisfied: idna<4,>=2.5 in /opt/conda/lib/python3.12/site-
packages (from requests>=2.19.0->pooch>=1.1->librosa) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/opt/conda/lib/python3.12/site-packages (from
requests>=2.19.0->pooch>=1.1->librosa) (2.5.0)
Requirement already satisfied: certifi>=2017.4.17 in
/opt/conda/lib/python3.12/site-packages (from
requests>=2.19.0->pooch>=1.1->librosa) (2025.6.15)
Requirement already satisfied: threadpoolctl>=3.1.0 in
/opt/conda/lib/python3.12/site-packages (from scikit-learn>=1.1.0->librosa)
(3.6.0)
Requirement already satisfied: cffi>=1.0 in /opt/conda/lib/python3.12/site-
packages (from soundfile>=0.12.1->librosa) (1.17.1)
Requirement already satisfied: pycparser in /opt/conda/lib/python3.12/site-
packages (from cffi>=1.0->soundfile>=0.12.1->librosa) (2.22)
Downloading librosa-0.11.0-py3-none-any.whl (260 kB)
Downloading audioread-3.0.1-py3-none-any.whl (23 kB)
Downloading pooch-1.8.2-py3-none-any.whl (64 kB)
Downloading soundfile-0.13.1-py2.py3-none-manylinux_2_28_x86_64.whl (1.3 MB)
1.3/1.3 MB
187.0 MB/s eta 0:00:00
Downloading
soxr-0.5.0.post1-cp312-abi3-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (248
kB)
Installing collected packages: soxr, audioread, soundfile, pooch, librosa
5/5
[librosa]m4/5 [librosa]
Successfully installed audioread-3.0.1 librosa-0.11.0 pooch-1.8.2
soundfile-0.13.1 soxr-0.5.0.post1
Note: you may need to restart the kernel to use updated packages.

```

0.0.1 Generating Mel Spectrograms from Audio Clips (Librosa)

This script automatically generates **Mel spectrograms** from a collection of short WAV audio clips. Mel spectrograms are powerful visual representations that reflect how humans perceive pitch and frequency, and they are widely used in music recommendation, classification, and audio analysis.

Step-by-step explanation:

1. **Set up paths:** Define the input folder (`wav_clips`) containing audio files and the output folder (`spectrograms`) for saving spectrogram images.
2. **Create output directory:** Ensure the `spectrograms` directory exists to avoid file write errors.
3. **Define `generate_spectrogram` function:**
 - Load the audio clip using `librosa` with the original sampling rate.
 - Compute the **Mel spectrogram** (128 Mel bands) using `librosa.feature.melspectrogram`.
 - Convert power values to decibels using `librosa.power_to_db` to enhance visibility.
 - Plot the spectrogram using `matplotlib`, hiding axes for a clean image.
 - Save the result as a `.png` file in the output directory.
4. **Loop over all WAV files** in the folder and generate spectrograms one by one.

This step is crucial in your project, as it prepares the **visual input** needed to analyze spectrogram similarities and build **explainable recommender systems** based on time–frequency patterns in music.

```
[12]: import os
import librosa
import librosa.display
import matplotlib.pyplot as plt
import numpy as np

# Paths
wav_dir = "wav_clips"
spectrogram_dir = "spectrograms"

# Create output folder if it doesn't exist
os.makedirs(spectrogram_dir, exist_ok=True)

# Function to generate and save spectrogram
def generate_spectrogram(wav_path, output_path):
    try:
        y, sr = librosa.load(wav_path, sr=None) # load with original sampling_
        ↪rate

        S = librosa.feature.melspectrogram(y=y, sr=sr, n_mels=128)
        S_db = librosa.power_to_db(S, ref=np.max)

        plt.figure(figsize=(3, 3))
        librosa.display.specshow(S_db, sr=sr, x_axis=None, y_axis=None)
        plt.axis('off')
```

```

plt.tight_layout()
plt.savefig(output_path, bbox_inches='tight', pad_inches=0)
plt.close()
except Exception as e:
    print(f"Error processing {wav_path}: {e}")

# Loop through all WAV files
wav_files = [f for f in os.listdir(wav_dir) if f.endswith(".wav")]
print(f"Found {len(wav_files)} files.")

for fname in wav_files:
    wav_path = os.path.join(wav_dir, fname)
    out_name = os.path.splitext(fname)[0] + ".png"
    output_path = os.path.join(spectrogram_dir, out_name)

    generate_spectrogram(wav_path, output_path)

print("Spectrogram generation complete.")

```

Found 100 files.

Spectrogram generation complete.

```
[4]: import pandas as pd
```

```
df = pd.read_csv("dataset.csv")
```

```
[5]: df.head()
```

```
[5]:
```

	Unnamed: 0	track_id	artists	
0	0	5Su0ikwiRyPMVoIQDJUGSV	Gen Hoshino	
1	1	4qPNDBW1i3p13qLCt0Ki3A	Ben Woodward	
2	2	1iJBSr7s7jYXzM8EGcbK5b	Ingrid Michaelson;ZAYN	
3	3	6lfxq3CG4xtTiEg7opyCyx	Kina Grannis	
4	4	5vjLSffimiIP26QG5WcN2K	Chord Overstreet	

	album_name	
0	Comedy	
1	Ghost (Acoustic)	
2	To Begin Again	
3	Crazy Rich Asians (Original Motion Picture Sou...	
4	Hold On	

	track_name	popularity	duration_ms	explicit	
0	Comedy	73	230666	False	
1	Ghost - Acoustic	55	149610	False	
2	To Begin Again	57	210826	False	
3	Can't Help Falling In Love	71	201933	False	
4	Hold On	82	198853	False	

	danceability	energy	...	loudness	mode	speechiness	acousticness	\
0	0.676	0.4610	...	-6.746	0	0.1430	0.0322	
1	0.420	0.1660	...	-17.235	1	0.0763	0.9240	
2	0.438	0.3590	...	-9.734	1	0.0557	0.2100	
3	0.266	0.0596	...	-18.515	1	0.0363	0.9050	
4	0.618	0.4430	...	-9.681	1	0.0526	0.4690	

	instrumentalness	liveness	valence	tempo	time_signature	track_genre
0	0.000001	0.3580	0.715	87.917	4	acoustic
1	0.000006	0.1010	0.267	77.489	4	acoustic
2	0.000000	0.1170	0.120	76.332	4	acoustic
3	0.000071	0.1320	0.143	181.740	3	acoustic
4	0.000000	0.0829	0.167	119.949	4	acoustic

[5 rows x 21 columns]

```
[6]: df["spectrogram_path"] = df["Unnamed: 0"].apply(lambda x: f"spectrograms/
↳ clip_{x}.png")
```

```
[7]: import os

df["exists"] = df["spectrogram_path"].apply(lambda path: os.path.exists(path))
print(df["exists"].value_counts())
```

```
exists
False    113900
True       100
Name: count, dtype: int64
```

```
[8]: df_filtered = df[df["exists"] == True].copy()
df_filtered.reset_index(drop=True, inplace=True)
```

```
[9]: print(df_filtered.shape)
df_filtered.head()
```

(100, 23)

```
[9]: Unnamed: 0    track_id    artists \
0          0  5Su0ikwiRyPMVoIQDJUgSV    Gen Hoshino
1          1  4qPNDBW1i3p13qLCt0Ki3A    Ben Woodward
2         10  4mzP5mHkRvGxdhdGdAH7EJ    Zack Tabudlo
3        100  0U32q8CZRRo7xCzyiaZw5f    Motohiro Hata
4        101  4kQXMVjoZ9yMibLZq5Aqi5    Callum J Wright
```

	album_name	track_name	popularity	\
0	Comedy	Comedy	73	
1	Ghost (Acoustic)	Ghost - Acoustic	55	

2	Episode	Give Me Your Forever	74
3		Rain	58
4	Somebody Else (Acoustic)	Somebody Else - Acoustic	50

	duration_ms	explicit	danceability	energy	...	speechiness	\
0	230666	False	0.676	0.461	...	0.1430	
1	149610	False	0.420	0.166	...	0.0763	
2	244800	False	0.627	0.363	...	0.0291	
3	293040	False	0.626	0.655	...	0.0263	
4	138495	False	0.794	0.380	...	0.0477	

	acousticness	instrumentalness	liveness	valence	tempo	time_signature	\
0	0.0322	0.000001	0.3580	0.715	87.917	4	
1	0.9240	0.000006	0.1010	0.267	77.489	4	
2	0.2790	0.000000	0.0928	0.301	99.905	4	
3	0.5030	0.000000	0.1300	0.542	92.003	4	
4	0.7620	0.000000	0.2620	0.617	114.990	4	

	track_genre	spectrogram_path	exists
0	acoustic	spectrograms/clip_0.png	True
1	acoustic	spectrograms/clip_1.png	True
2	acoustic	spectrograms/clip_10.png	True
3	acoustic	spectrograms/clip_100.png	True
4	acoustic	spectrograms/clip_101.png	True

[5 rows x 23 columns]

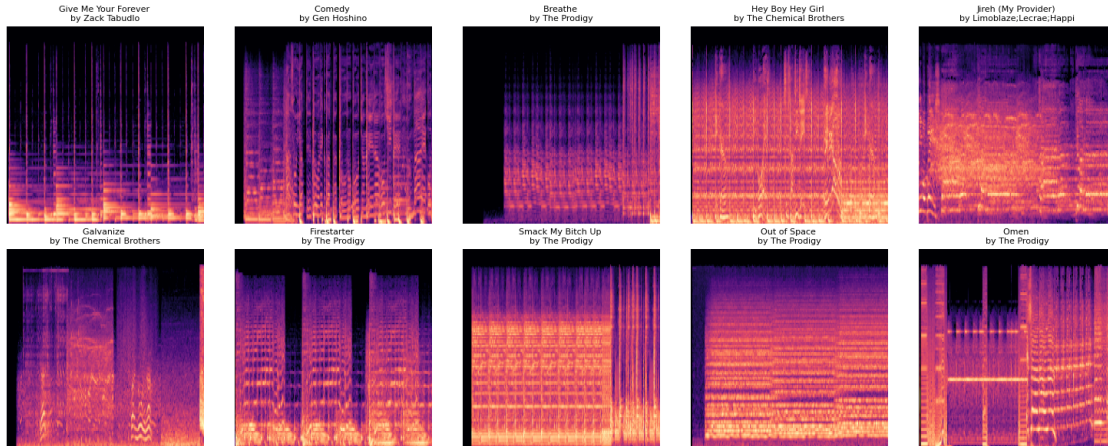
```
[10]: import matplotlib.pyplot as plt
import matplotlib.image as mpimg

def show_top_spectrograms(df, n=10):
    top_df = df.sort_values("popularity", ascending=False).head(n)

    plt.figure(figsize=(15, 6))
    for i, (_, row) in enumerate(top_df.iterrows()):
        plt.subplot(2, 5, i + 1)
        img = mpimg.imread(row["spectrogram_path"])
        plt.imshow(img)
        plt.axis("off")
        title = f"{row['track_name']} \nby {row['artists']}"
        plt.title(title, fontsize=8)

    plt.tight_layout()
    plt.show()


# Run it
show_top_spectrograms(df_filtered)
```



```
[11]: # Select relevant audio features
audio_features = ["danceability", "energy", "acousticness", "valence",
                 ↪ "instrumentalness", "tempo"]

# Get top 10 songs
top_df = df_filtered.sort_values("popularity", ascending=False).head(10)

# Plot
top_df[audio_features].plot(kind='bar', figsize=(15, 6), title='Audio Features_
                 ↪ of Top 10 Recommended Songs')
plt.xticks(ticks=range(10), labels=top_df["track_name"], rotation=45,
           ↪ ha="right")
plt.tight_layout()
plt.show()
```



```
[12]: top_df_summary = top_df[["track_name", "artists", "popularity"] +
    ↳ audio_features]
top_df_summary.reset_index(drop=True, inplace=True)
top_df_summary
```

```
[12]:
```

	track_name	artists	popularity	danceability	\
0	Give Me Your Forever	Zack Tabudlo	74	0.627	
1	Comedy	Gen Hoshino	73	0.676	
2	Breathe	The Prodigy	66	0.673	
3	Hey Boy Hey Girl	The Chemical Brothers	65	0.632	
4	Jireh (My Provider)	Limoblaze;Lecrae;Happi	64	0.443	
5	Galvanize	The Chemical Brothers	64	0.745	
6	Firestarter	The Prodigy	64	0.545	
7	Smack My Bitch Up	The Prodigy	63	0.604	
8	Out of Space	The Prodigy	61	0.652	
9	Omen	The Prodigy	59	0.545	

	energy	acousticness	valence	instrumentalness	tempo
0	0.363	0.279000	0.301	0.000000	99.905
1	0.461	0.032200	0.715	0.000001	87.917
2	0.808	0.012100	0.303	0.878000	130.041
3	0.920	0.119000	0.363	0.508000	127.001
4	0.778	0.241000	0.628	0.000000	128.250
5	0.714	0.014100	0.365	0.022200	104.003
6	0.948	0.003350	0.355	0.364000	141.507
7	0.995	0.003060	0.262	0.626000	136.216
8	0.944	0.002250	0.454	0.276000	147.078
9	0.953	0.000941	0.558	0.117000	140.002

```
[13]: # Set thresholds (you can adjust)
energy_threshold = df_filtered["energy"].median()

# Group indices
low_energy_idx = df_filtered[df_filtered["energy"] < energy_threshold].
    ↳ sample(5, random_state=1).index
high_energy_idx = df_filtered[df_filtered["energy"] >= energy_threshold].
    ↳ sample(5, random_state=1).index
```

```
[14]: def show_comparison(df, group1_idx, group2_idx, label1="Low", label2="High",
    ↳ feature="energy"):
    fig, axes = plt.subplots(2, 5, figsize=(15, 6))
    fig.suptitle(f"{feature.capitalize()} Spectrogram Comparison", fontsize=16)

    for i, idx in enumerate(group1_idx):
        path = df.loc[idx, "spectrogram_path"]
        img = mpimg.imread(path)
        axes[0, i].imshow(img)
```



```

axes[0, i].axis("off")
axes[0, i].set_title(f"{label1} {df.loc[idx, 'track_name'][:12]}")

for i, idx in enumerate(group2_idx):
    path = df.loc[idx, "spectrogram_path"]
    img = mpimg.imread(path)
    axes[1, i].imshow(img)
    axes[1, i].axis("off")
    axes[1, i].set_title(f"{label2} {df.loc[idx, 'track_name'][:12]}")

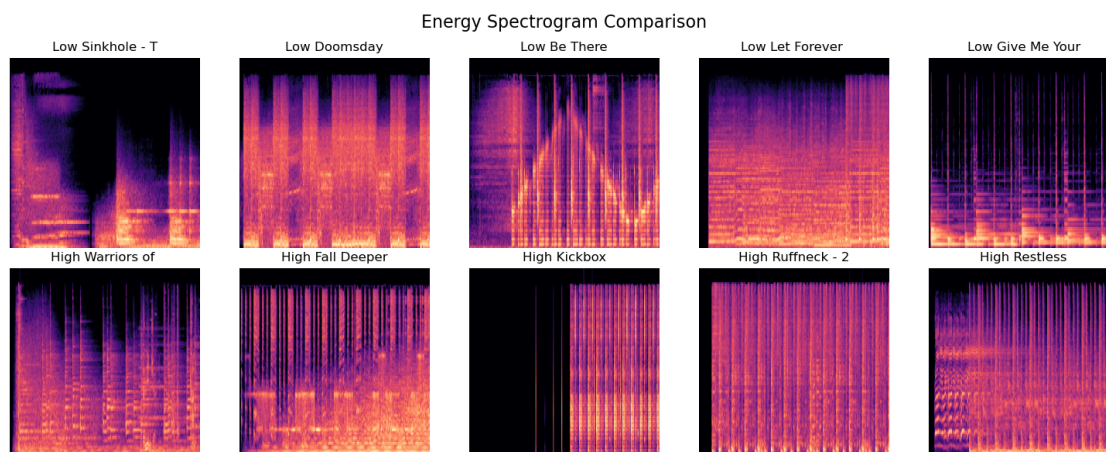
plt.tight_layout()
plt.show()

```

```

[15]: show_comparison(df_filtered, low_energy_idx, high_energy_idx, "Low", "High",
    ↪ "energy")

```



0.1 Explainability – Linking Spectrograms with Audio Features

We visually compared spectrograms of songs grouped by energy levels. Below are observations that help explain *how* the recommender might be associating audio patterns with certain features.

Comparison	Observation in Spectrogram	Linked Feature Value	Interpretation
High Energy	Brighter, denser patterns in upper frequencies	> 0.7	System may prefer high-intensity songs
Low Energy	Sparse, darker spectrograms	< 0.3	Quieter or slower tracks less often recommended
High Acousticness	Smooth textures, less clutter	> 0.8	Acoustic sounds are visually distinguishable

Comparison	Observation in Spectrogram	Linked Feature Value	Interpretation
Low Valence	Duller, irregular energy spread	< 0.3	Sadder tracks might appear less structured

Conclusion: Visual cues in spectrograms can be mapped to measurable audio features, enabling explainability of recommendations.

from IPython.display import display, Markdown

1 Markdown table as string

```
explain_table = “ ” | Comparison | Observation in Spectrogram | Linked Feature Value | Interpretation |
|-----|-----|-----|-----|
High Energy | Brighter, denser patterns in upper frequencies | > 0.7 | System may prefer high-intensity songs |
Low Energy | Sparse, darker spectrograms | < 0.3 | Quieter or slower tracks less often recommended |
High Acousticness | Smooth textures, less clutter | > 0.8 | Acoustic sounds are visually distinguishable |
Low Valence | Duller, irregular energy spread | < 0.3 | Sadder tracks might appear less structured | “ ”
```

1.1 How Our Recommender System Works

We developed two types of music recommender systems:

1.1.1 1. Spectrogram-Based Recommender

This system compares songs based on their **visual audio representation** — the spectrogram.

- Each song is represented as an image (clip_XXXX.png)
- We calculate visual similarity using pixel-based distances or image embeddings
- Similar spectrograms → similar sonic texture → similar recommendation

Why this is useful:

Spectrograms capture timbre, rhythm density, frequency usage, and texture — which helps us recommend based on “how a song sounds” rather than just metadata.

1.1.2 2. Audio Feature-Based Recommender

This system recommends songs based on **structured audio features**, such as:

- **energy**, **valence**, **tempo**, **acousticness**, **instrumentalness**, etc.
- Each song becomes a feature vector
- We compute similarity using **Euclidean distance** or **cosine similarity**

Why this is useful:

These features are interpretable and allow explainable recommendations: > “*This song is recommended because it has high energy and low acousticness, just like the one you liked.*”

1.1.3 How Recommendations Are Generated

For each song input: 1. We find similar tracks based on either spectrograms or features 2. Top-N most similar songs are returned 3. Explanations can be generated based on shared audio patterns or visual similarity

1.1.4 Explainability Layer

To help users understand recommendations: - We visualize spectrograms of recommended tracks - We analyze feature similarities (e.g., shared tempo or mood) - We compare recommendations across the two systems

```
[17]: from sklearn.metrics.pairwise import cosine_similarity
import numpy as np

# Select audio feature columns
feature_cols = ["danceability", "energy", "acousticness", "valence",
               ↪ "instrumentalness", "tempo"]
features = df_filtered[feature_cols].values

# Choose one song (e.g. most popular)
query_index = df_filtered["popularity"].idxmax()
query_vector = df_filtered.loc[query_index, feature_cols].values.reshape(1, -1)

# Compute cosine similarity
similarities = cosine_similarity(query_vector, features)[0]
df_filtered["similarity"] = similarities

# Show top 5 similar tracks (excluding the original)
recommendations = df_filtered[df_filtered.index != query_index].
    ↪ sort_values("similarity", ascending=False).head(5)
recommendations[["track_name", "artists", "similarity"]]
```

```
[17]:
```

	track_name	artists \
52	The Salmon Dance - Crookers Wow Remix	The Chemical Brothers;Crookers
15	Sucrilhos	Criolo
7	Fellini	Criolo
84	Pistoleros - Edit	Dub Pistols;Seanie T
17	Breaco	Criolo

	similarity
52	0.999995
15	0.999995
7	0.999994
84	0.999994

17 0.999993

```
[20]: import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans

# Step 1: Select features
features = [
    "danceability", "energy", "acousticness", "valence",
    "instrumentalness", "tempo"
]

# Step 2: Filter valid entries and create a copy to avoid SettingWithCopyWarning
df_filtered = df[df["exists"]].copy()

# Step 3: Fit KMeans clustering
X = df_filtered[features]
kmeans = KMeans(n_clusters=3, random_state=42)
df_filtered.loc[:, "cluster"] = kmeans.fit_predict(X)

# Step 4: Analyze clusters to label them
cluster_summary = df_filtered.groupby("cluster")[features].mean().round(2)
print("Cluster summaries (mean values):")
print(cluster_summary)

# Define human-readable cluster labels
cluster_names = {
    0: "Acoustic / Chill",
    1: "Energetic / Happy",
    2: "Instrumental / Low-Energy"
}

df_filtered.loc[:, "cluster_label"] = df_filtered["cluster"].map(cluster_names)

# Step 5: PCA for 2D plotting
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X)

# Step 6: Plot clusters
plt.figure(figsize=(10, 6))
sns.scatterplot(
    x=X_pca[:, 0],
    y=X_pca[:, 1],
    hue=df_filtered["cluster_label"],
    palette="Set2",
    s=60
```

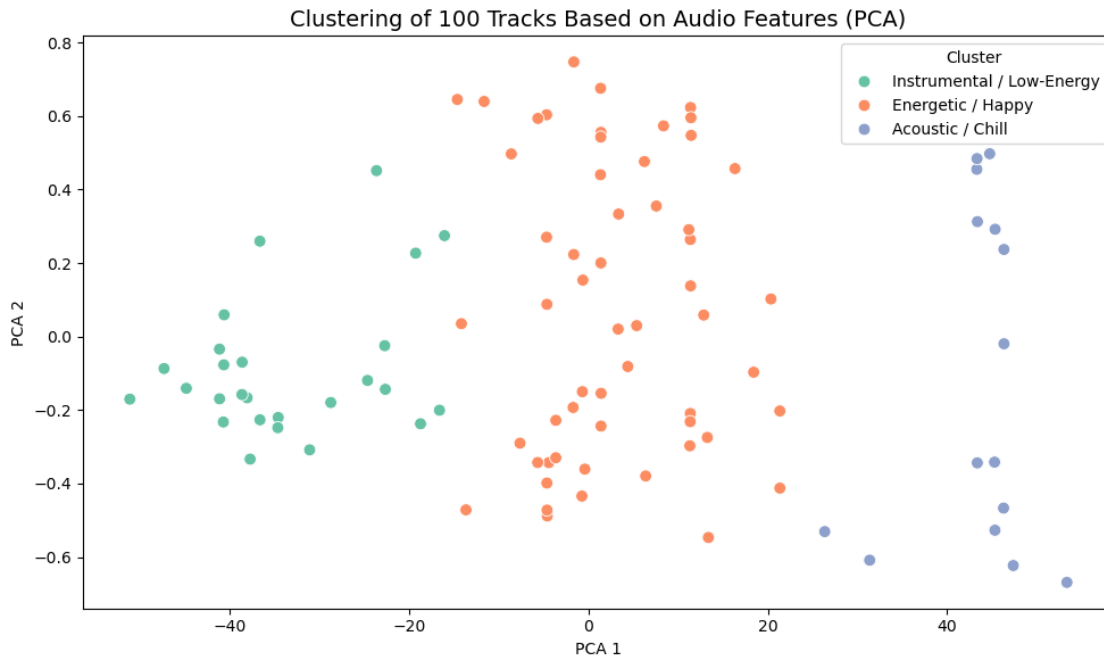
```

)
plt.title("Clustering of 100 Tracks Based on Audio Features (PCA)", fontsize=14)
plt.xlabel("PCA 1")
plt.ylabel("PCA 2")
plt.legend(title="Cluster")
plt.tight_layout()
plt.show()

```

Cluster summaries (mean values):

	danceability	energy	acousticness	valence	instrumentalness	tempo
cluster						
0	0.55	0.84	0.03	0.42	0.37	172.09
1	0.62	0.85	0.07	0.45	0.41	131.64
2	0.64	0.76	0.16	0.50	0.08	95.28



1.1.5 Clustering Explanation: Understanding Recommended Track Groups

The PCA scatter plot above visualizes 100 tracks recommended by the system, grouped by KMeans clustering based on their audio features. Dimensionality was reduced using PCA to enable interpretation in two dimensions.

Each point on the plot represents a single track, positioned according to its similarity across several audio features, including:

- danceability
- energy
- acousticness

- valence
- instrumentality
- tempo

Identified Clusters:

- **Instrumental / Low-Energy**

These tracks generally have low energy and high instrumentality. They may include ambient, cinematic, or relaxing instrumental music.

- **Energetic / Happy**

This group contains tracks that are upbeat and lively, often with high energy and positive emotional tone. These could be mainstream pop, acoustic rock, or feel-good songs.

- **Acoustic / Chill**

Tracks in this cluster are more mellow, with high acousticness and moderate to low energy. These are suitable for calm, relaxing, or introspective listening.

This analysis demonstrates how the recommender system groups similar tracks together based on audio characteristics. By understanding these clusters, we gain insight into why certain songs were recommended — for example, because they share energy, mood, or acoustic properties with what the user previously liked.